

Neural Networks:

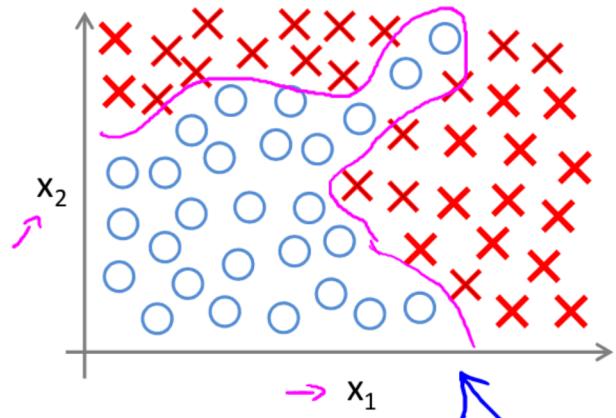
Representation

Non-linear

hypotheses

Machine Learning

Non-linear Classification



- x_1 = size
- x_2 = # bedrooms
- x_3 = # floors
- x_4 = age
- ...
- x_{100} -

$\left. \begin{array}{c} \\ \\ \\ \\ \end{array} \right\} h=100$

$$g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1 x_2 + \theta_4 x_1^2 x_2 + \theta_5 x_1^3 x_2 + \theta_6 x_1^2 x_2^2 + \dots)$$

$$\rightarrow x_1^2, x_1 x_2, x_1 x_3, x_1 x_4 \dots x_1 x_{100}$$

$$x_2^2, x_2 x_3 \dots$$

≈ 5000 feature

$O(n^2)$

$$\rightarrow \underline{x_1^2}, \underline{x_2^2}, \underline{x_3^2}, \dots, \underline{x_{100}^2}$$

$$\rightarrow \underline{x_1 x_2 x_3}, \underline{x_1^2 x_2}, \underline{x_{10} x_{11} x_{17}}, \dots$$

$O(n^3)$

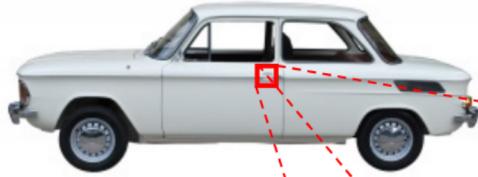
$170,000$

$\frac{n^2}{2}$

Andrew Ng

What is this?

You see this:



But the camera sees this:

194	210	201	212	199	213	215	195	178	158	182	209
180	189	190	221	209	205	191	167	147	115	129	163
114	126	140	188	176	165	152	140	170	106	78	88
87	103	115	154	143	142	149	153	173	101	57	57
102	112	106	131	122	138	152	147	128	84	58	66
94	95	79	104	105	124	129	113	107	87	69	67
68	71	69	98	89	92	98	95	89	88	76	67
41	56	68	99	63	45	60	82	58	76	75	65
20	43	69	75	56	41	51	73	55	70	63	44
50	50	57	69	75	75	73	74	53	68	59	37
72	59	53	66	84	92	84	74	57	72	63	42
67	61	58	65	75	78	76	73	59	75	69	50



Andrew Ng

Computer Vision: Car detection



Cars



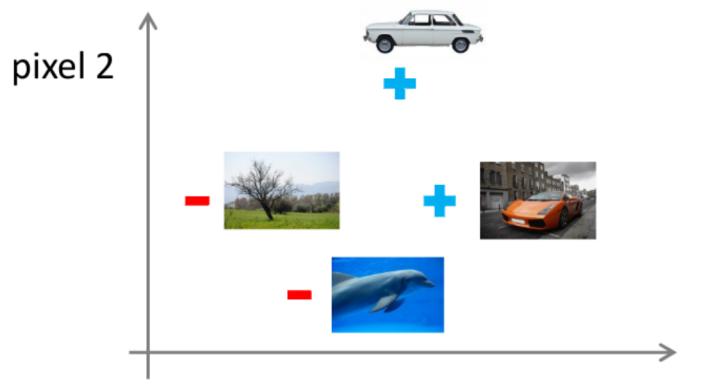
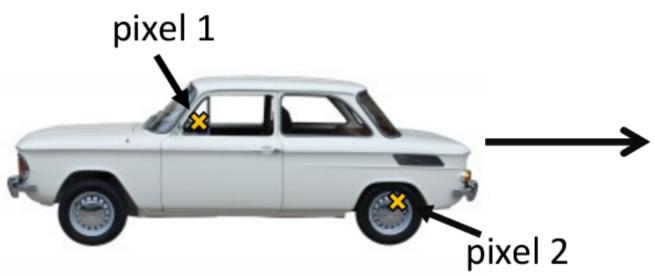
Not a car

Testing:



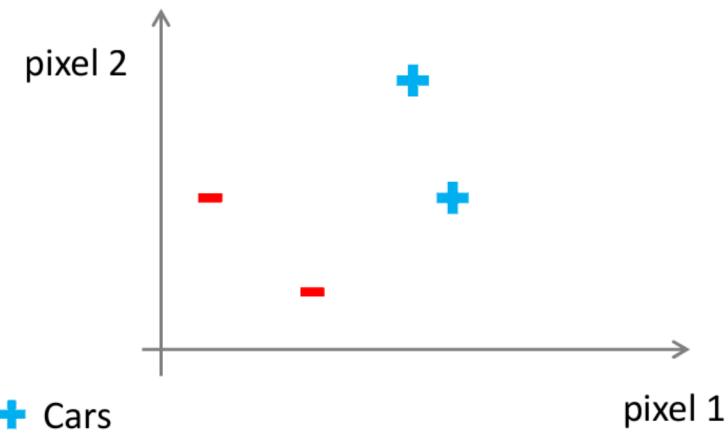
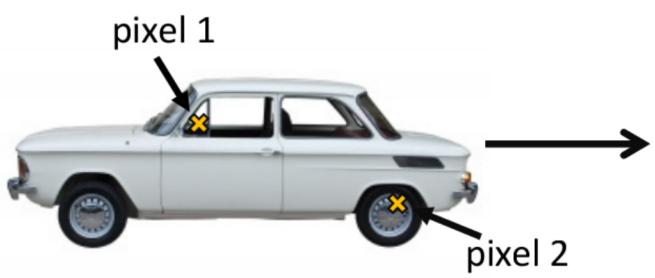
What is this?

Andrew Ng

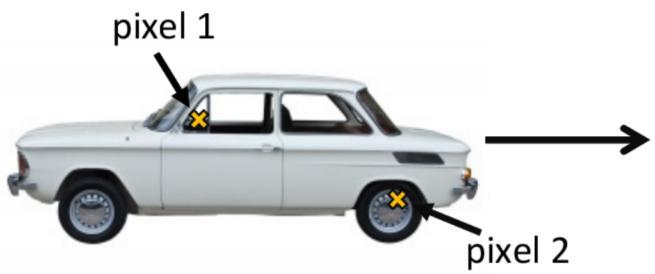


⊕ Cars
⊖ “Non”-Cars

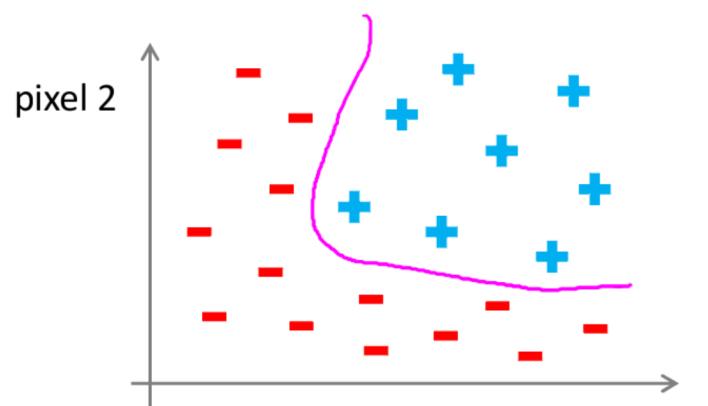
Andrew Ng



Andrew Ng



Learning
Algorithm



50×50 pixel images \rightarrow 2500 pixels

$n = 2500$ (7500 if RGB)

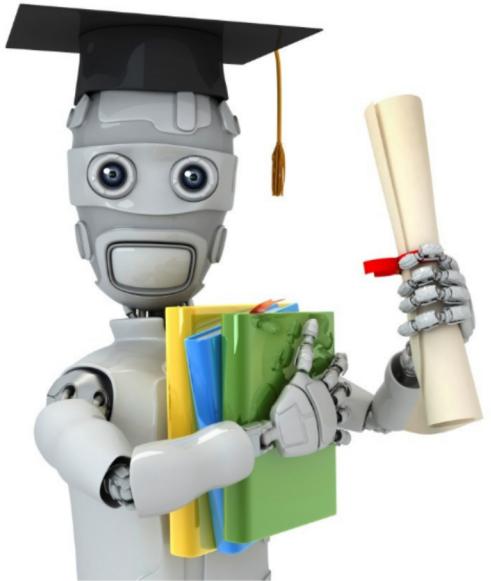
$$\rightarrow x = \begin{bmatrix} \text{pixel 1 intensity} \\ \text{pixel 2 intensity} \\ \vdots \\ \text{pixel 2500 intensity} \end{bmatrix}$$

$0-255$

- ⊕ Cars
- "Non"-Cars

Quadratic features ($x_i \times x_j$): ≈ 3 million features

Andrew Ng



Neural Networks: Representation

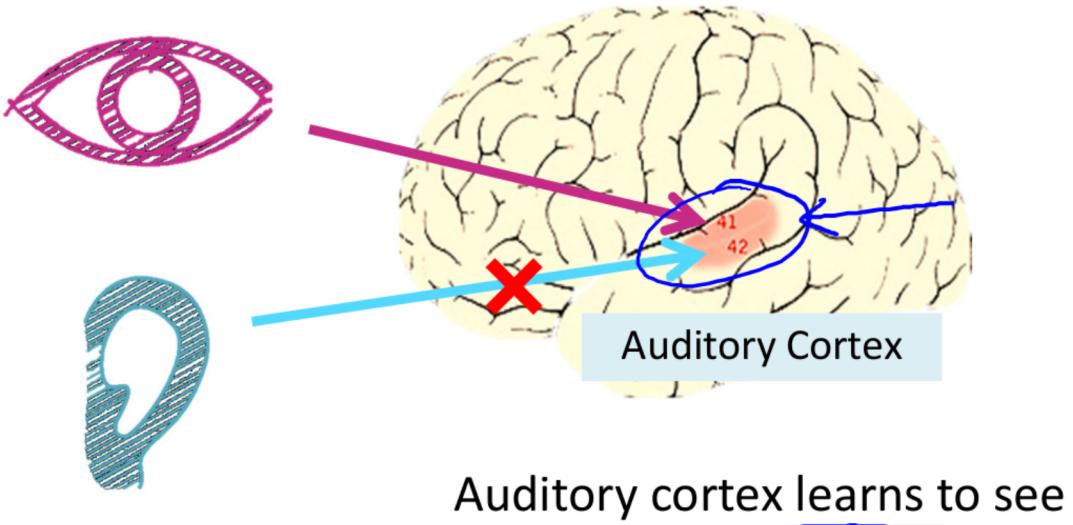
Neurons and the brain

Machine Learning

Neural Networks

- Origins: Algorithms that try to mimic the brain.
- Was very widely used in 80s and early 90s; popularity diminished in late 90s.
- Recent resurgence: State-of-the-art technique for many applications

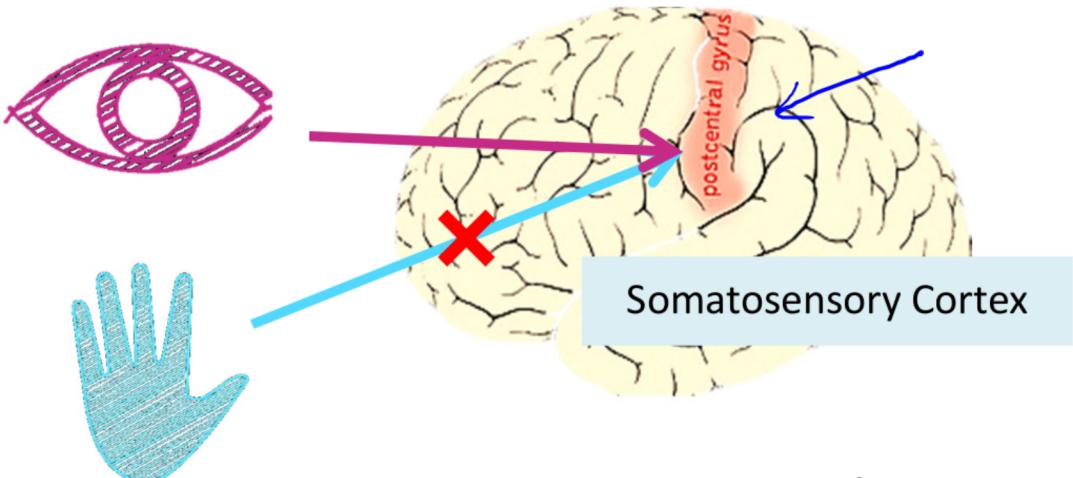
The “one learning algorithm” hypothesis



[Roe et al., 1992]

Andrew Ng

The “one learning algorithm” hypothesis



Somatosensory cortex learns to see

[Metin & Frost, 1989]

Andrew Ng

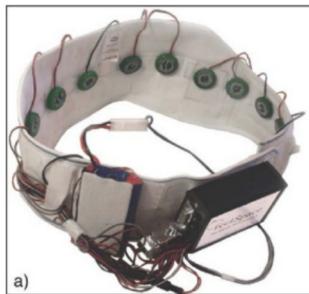
Sensor representations in the brain



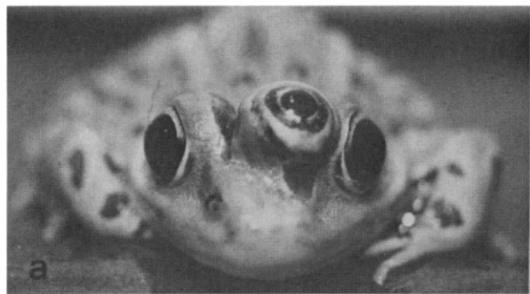
Seeing with your tongue



Human echolocation (sonar)



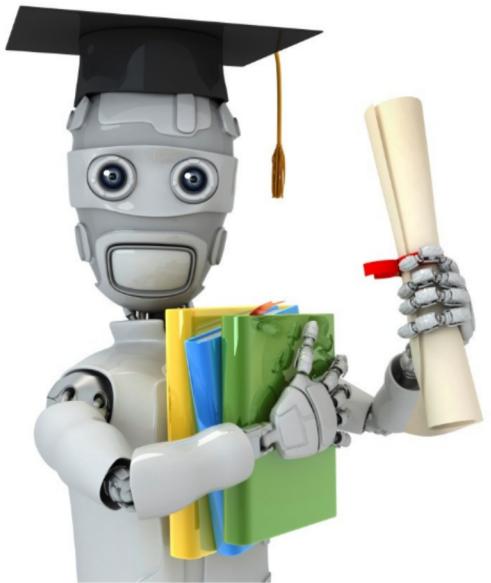
Haptic belt: Direction sense



Implanting a 3rd eye

[BrainPort; Welsh & Blasch, 1997; Nagel et al., 2005; Constantine-Paton & Law, 2009]

Andrew Ng



Neural Networks:

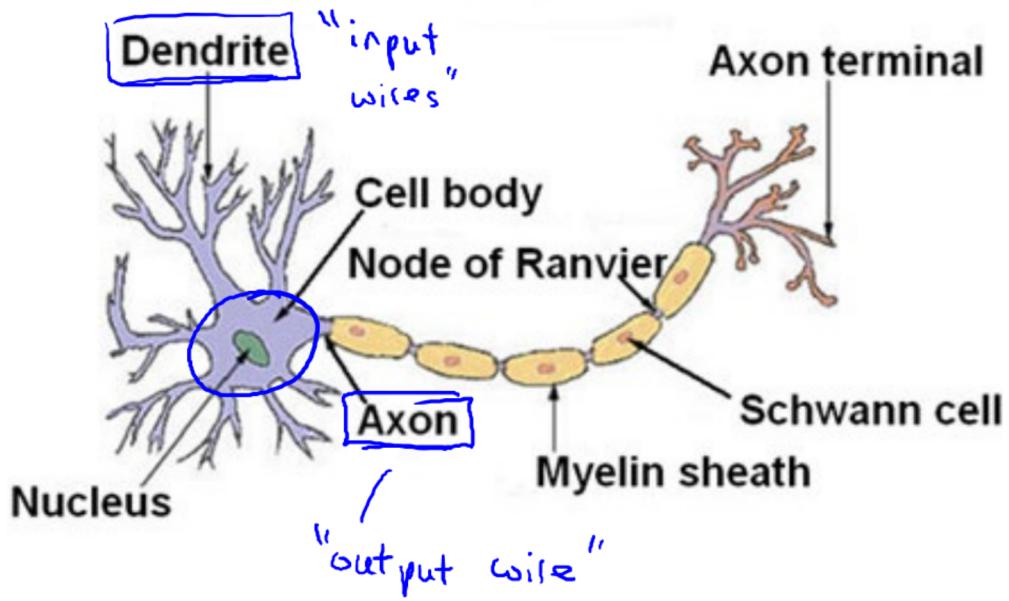
Representation

Model

representation I

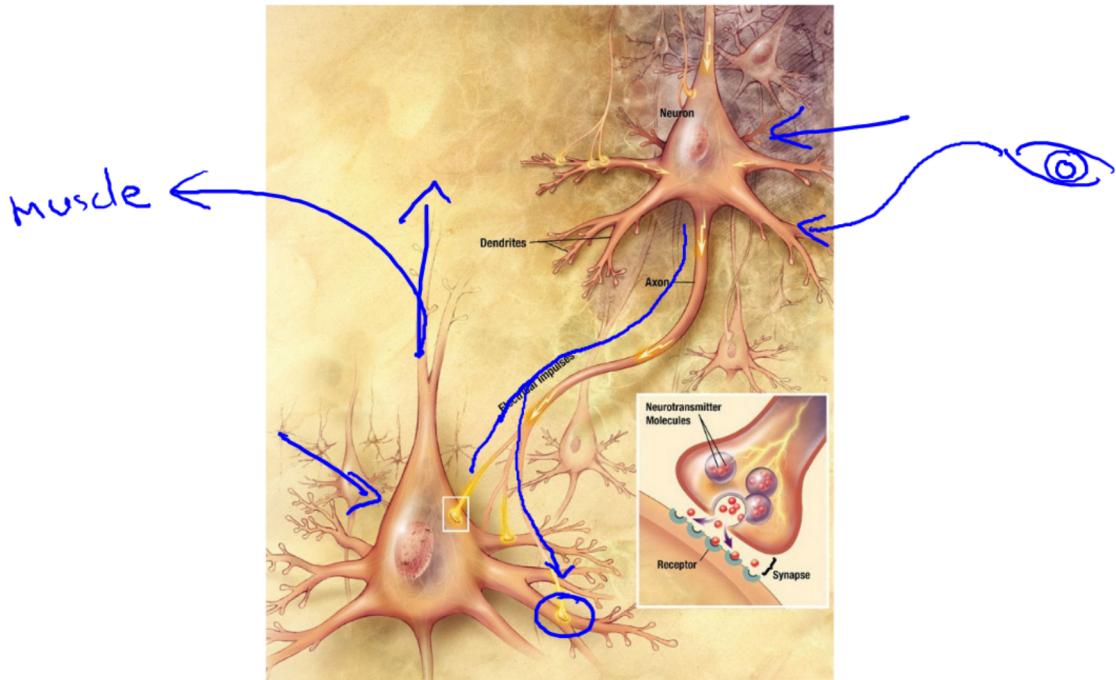
Machine Learning

Neuron in the brain



Andrew Ng

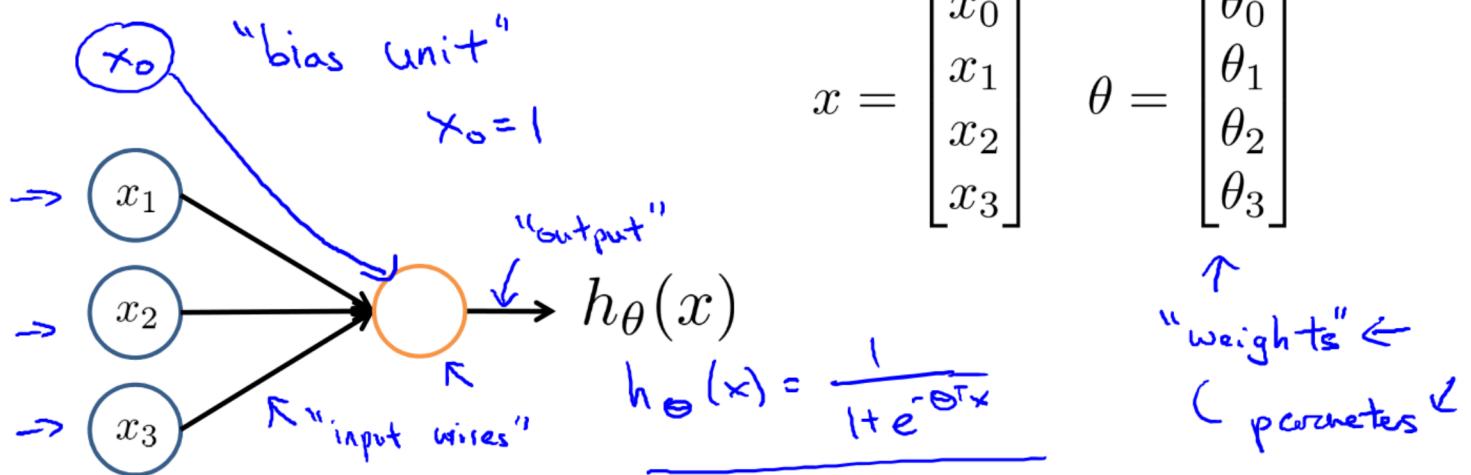
Neurons in the brain



[Credit: US National Institutes of Health, National Institute on Aging]

Andrew Ng

Neuron model: Logistic unit

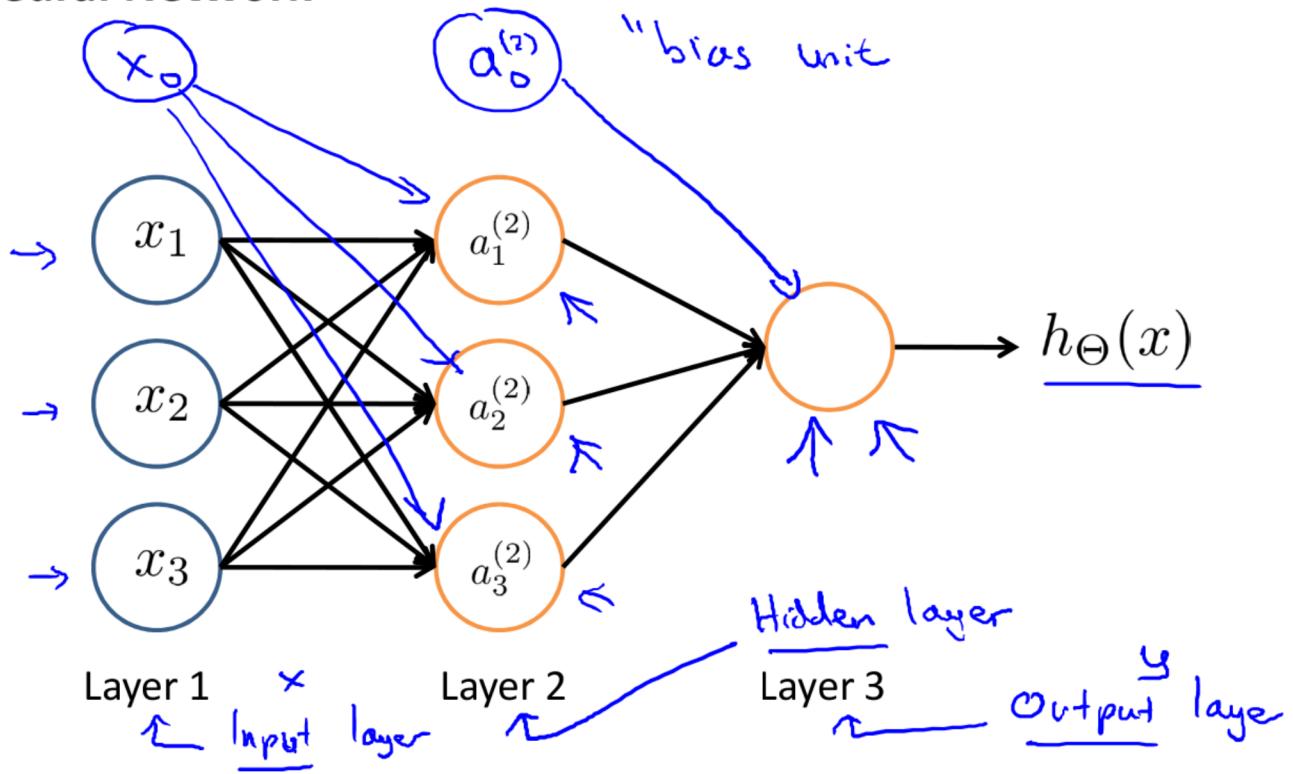


Sigmoid (logistic) activation function.

$$g(z) = \frac{1}{1+e^{-z}}$$

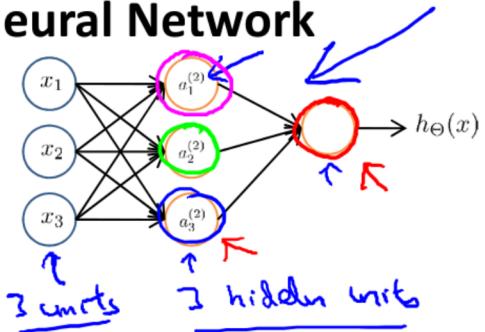
Andrew Ng

Neural Network



Andrew Ng

Neural Network



$\rightarrow a_i^{(j)}$ = "activation" of unit i in layer j

$\rightarrow \Theta^{(j)}$ = matrix of weights controlling function mapping from layer j to layer $j + 1$

$$h_\Theta(x)$$

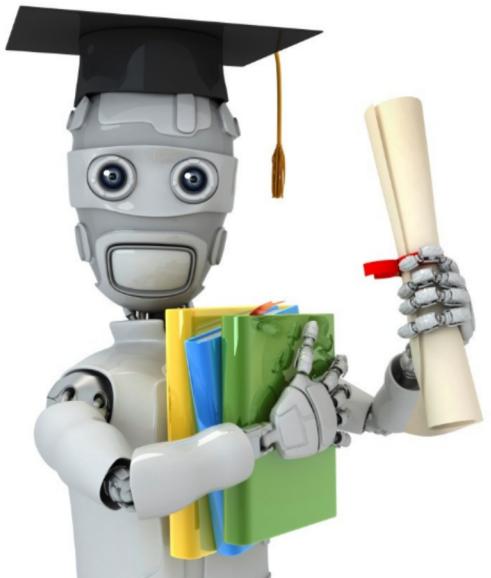
$$\rightarrow a_1^{(2)} = g(\Theta_{10}^{(1)}x_0 + \Theta_{11}^{(1)}x_1 + \Theta_{12}^{(1)}x_2 + \Theta_{13}^{(1)}x_3)$$

$$\rightarrow a_2^{(2)} = g(\Theta_{20}^{(1)}x_0 + \Theta_{21}^{(1)}x_1 + \Theta_{22}^{(1)}x_2 + \Theta_{23}^{(1)}x_3)$$

$$\rightarrow a_3^{(2)} = g(\Theta_{30}^{(1)}x_0 + \Theta_{31}^{(1)}x_1 + \Theta_{32}^{(1)}x_2 + \Theta_{33}^{(1)}x_3)$$

$$\rightarrow h_\Theta(x) = a_1^{(3)} = g(\Theta_{10}^{(2)}a_0^{(2)} + \Theta_{11}^{(2)}a_1^{(2)} + \Theta_{12}^{(2)}a_2^{(2)} + \Theta_{13}^{(2)}a_3^{(2)})$$

→ If network has s_j units in layer j , s_{j+1} units in layer $j + 1$, then $\Theta^{(j)}$ will be of dimension $s_{j+1} \times (s_j + 1)$. $\hookrightarrow s_{j+1} \times (s_j + 1)$



Neural Networks:

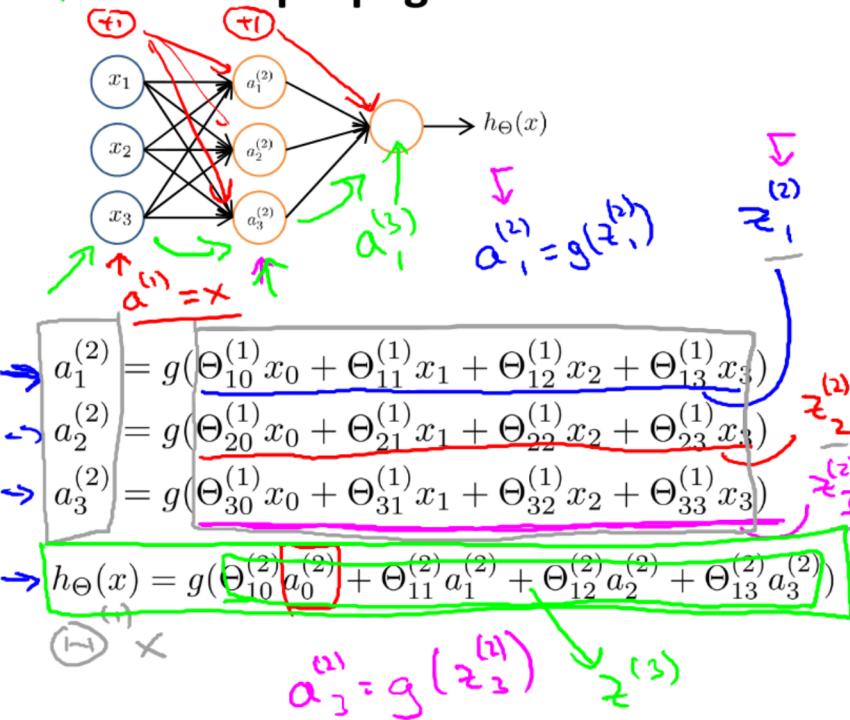
Representation

Model

representation II

Machine Learning

Forward propagation: Vectorized implementation



$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \end{bmatrix} \quad z^{(2)} = \begin{bmatrix} z_1^{(2)} \\ z_2^{(2)} \\ z_3^{(2)} \end{bmatrix}$$

$$z^{(2)} = \Theta^{(1)} \times a^{(1)}$$

$$a^{(2)} = g(z^{(2)})$$

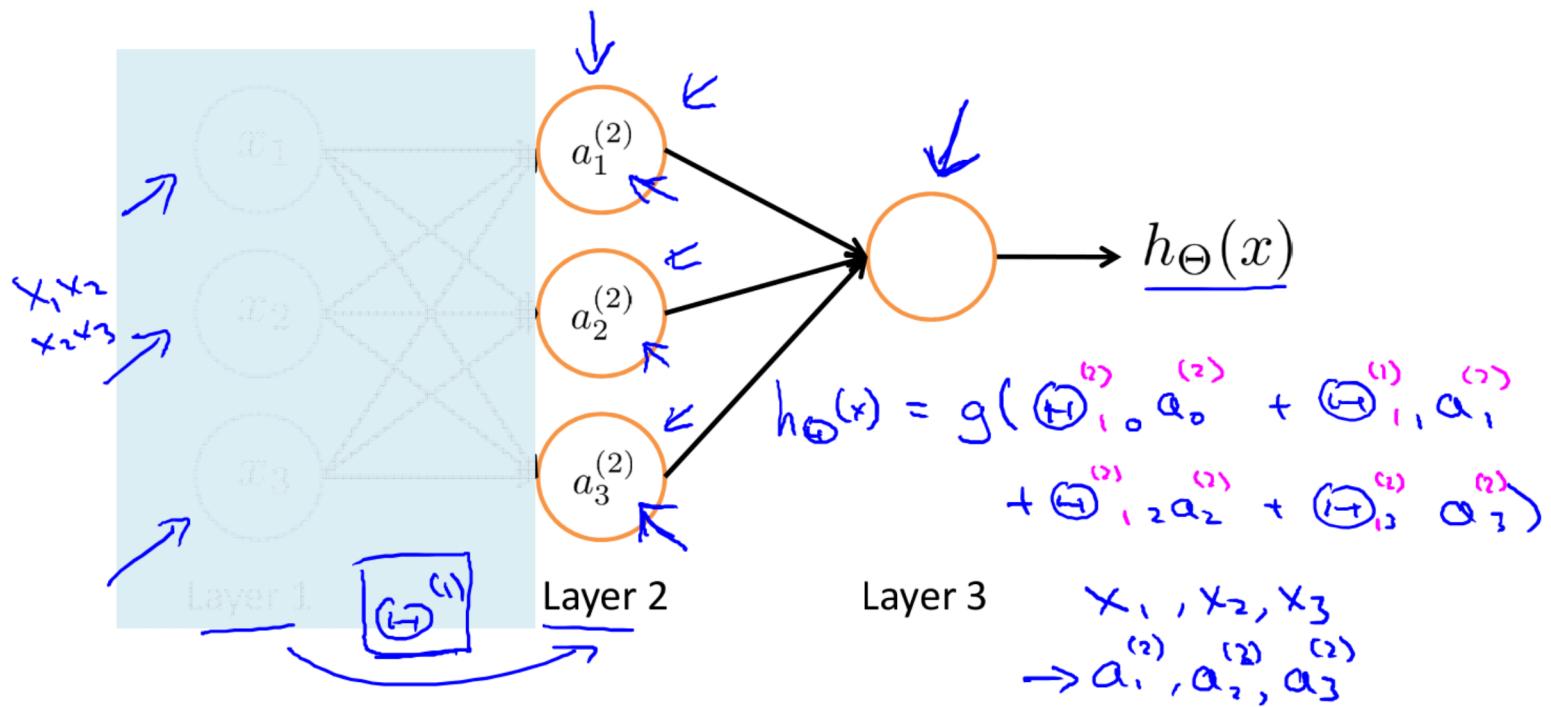
Add $a_0^{(2)} = 1.$ $\rightarrow a^{(2)} \in \mathbb{R}^4$

$$z^{(3)} = \Theta^{(2)} a^{(2)}$$

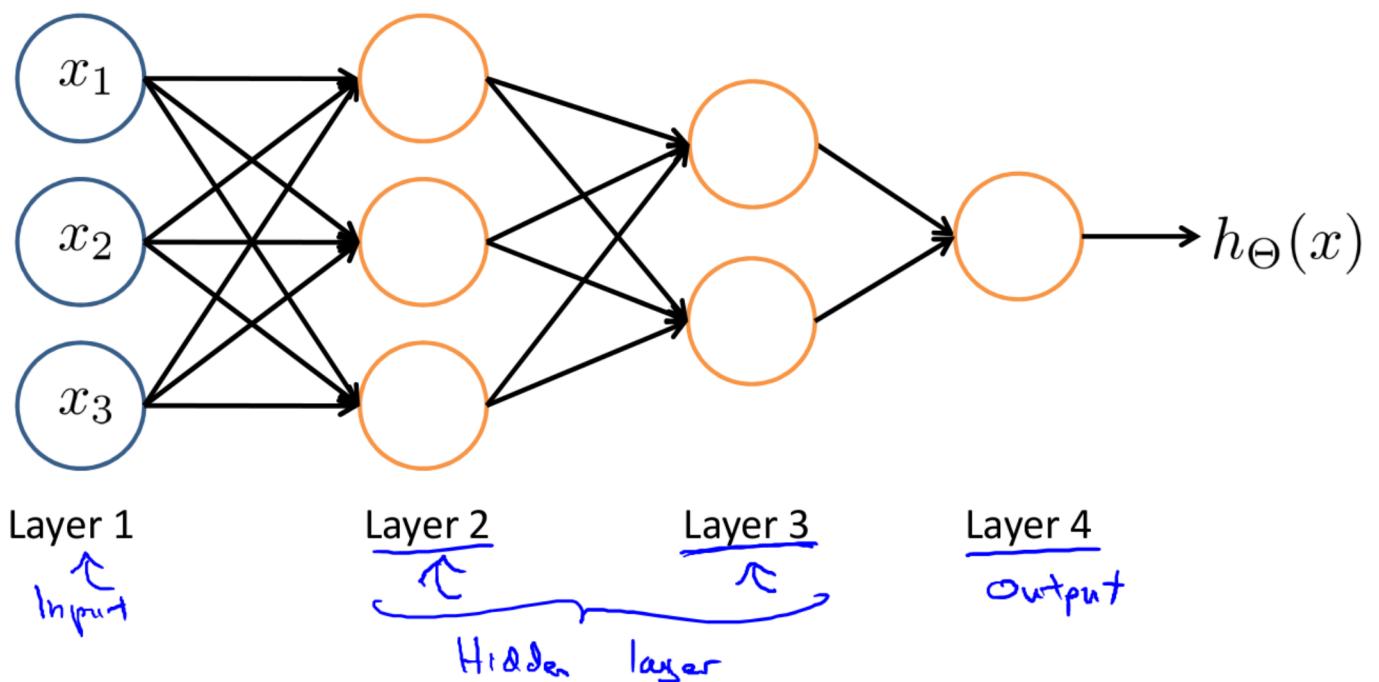
$$h_{\Theta}(x) = a^{(3)} = g(z^{(3)})$$

Andrew Ng

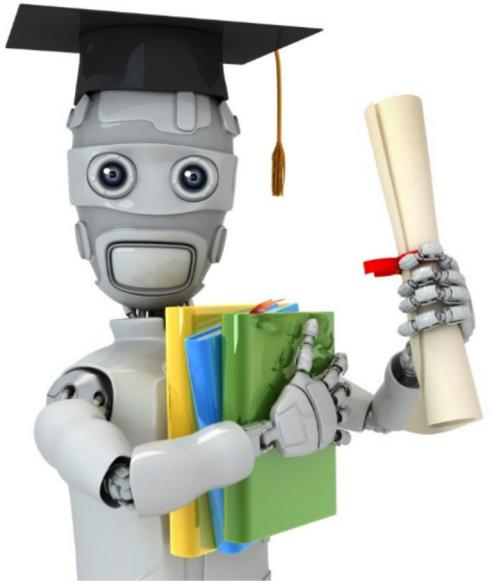
Neural Network learning its own features



Other network architectures



Andrew Ng



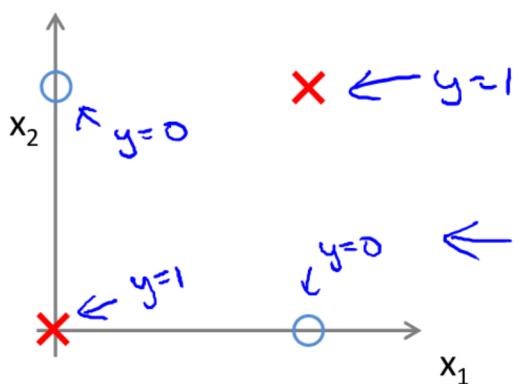
Neural Networks: Representation

Examples and intuitions I

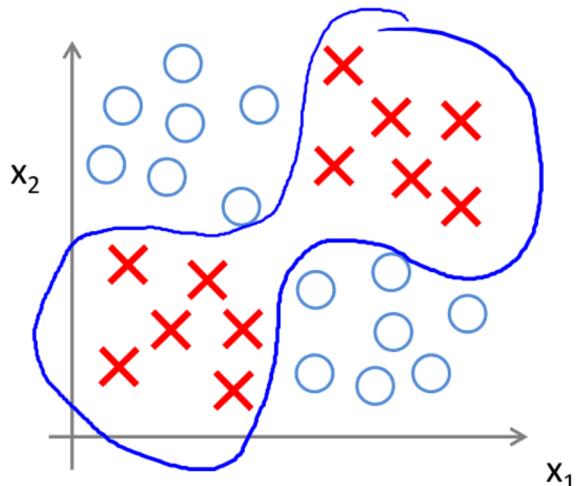
Machine Learning

Non-linear classification example: XOR/XNOR

→ x_1, x_2 are binary (0 or 1).



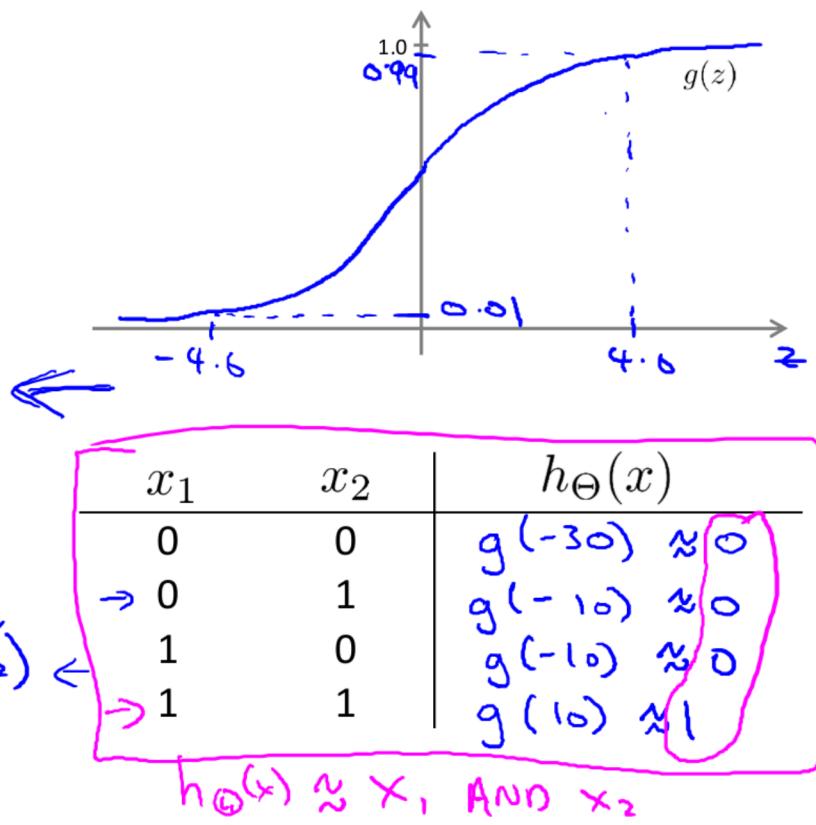
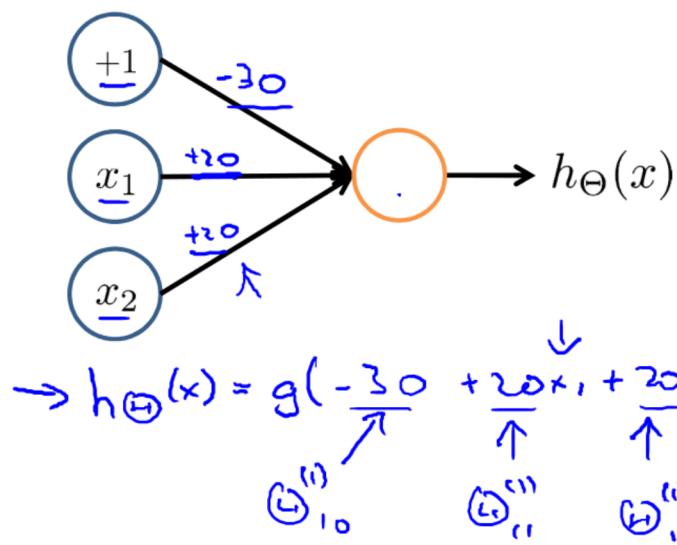
$$y = \underbrace{x_1 \text{ XOR } x_2}_{\substack{\rightarrow \\ \text{XNOR}}} \leftarrow$$
$$\underbrace{\text{NOT} (x_1 \text{ XOR } x_2)}_{\substack{\rightarrow \\ \text{XNOR}}}$$



Andrew Ng

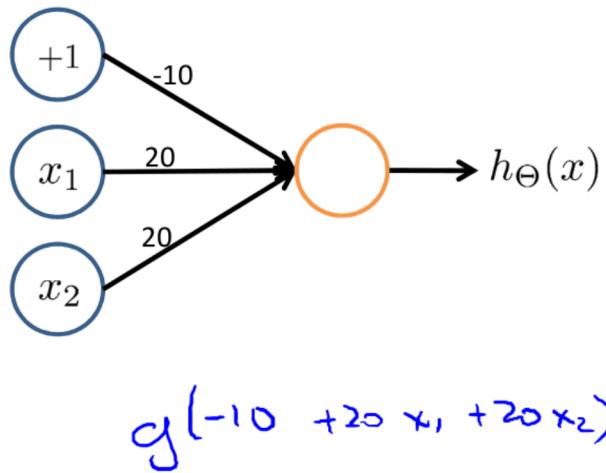
Simple example: AND

- $x_1, x_2 \in \{0, 1\}$
- $y = x_1 \text{ AND } x_2$



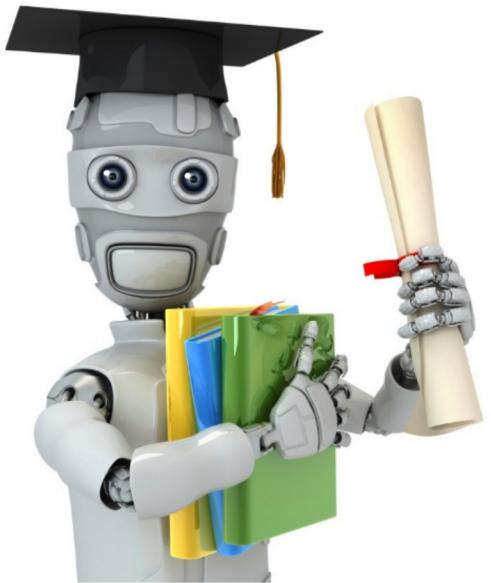
Andrew Ng

Example: OR function



x_1	x_2	$h_{\Theta}(x)$
0	0	$g(-10) \approx 0$
0	1	$g(10) \approx 1$
1	0	≈ 1
1	1	≈ 1

Andrew Ng



Neural Networks: Representation

Examples and intuitions II

Machine Learning

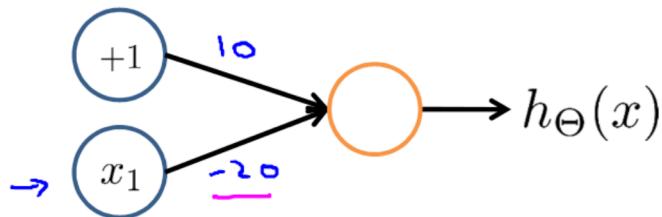
$\rightarrow x_1 \text{ AND } x_2$

$\rightarrow x_1 \text{ OR } x_2$

$\{0, 1\}$.

Negation:

NOT x_1



x_1	$h_{\Theta}(x)$
0	$g(10) \approx 1$
1	$g(-20) \approx 0$

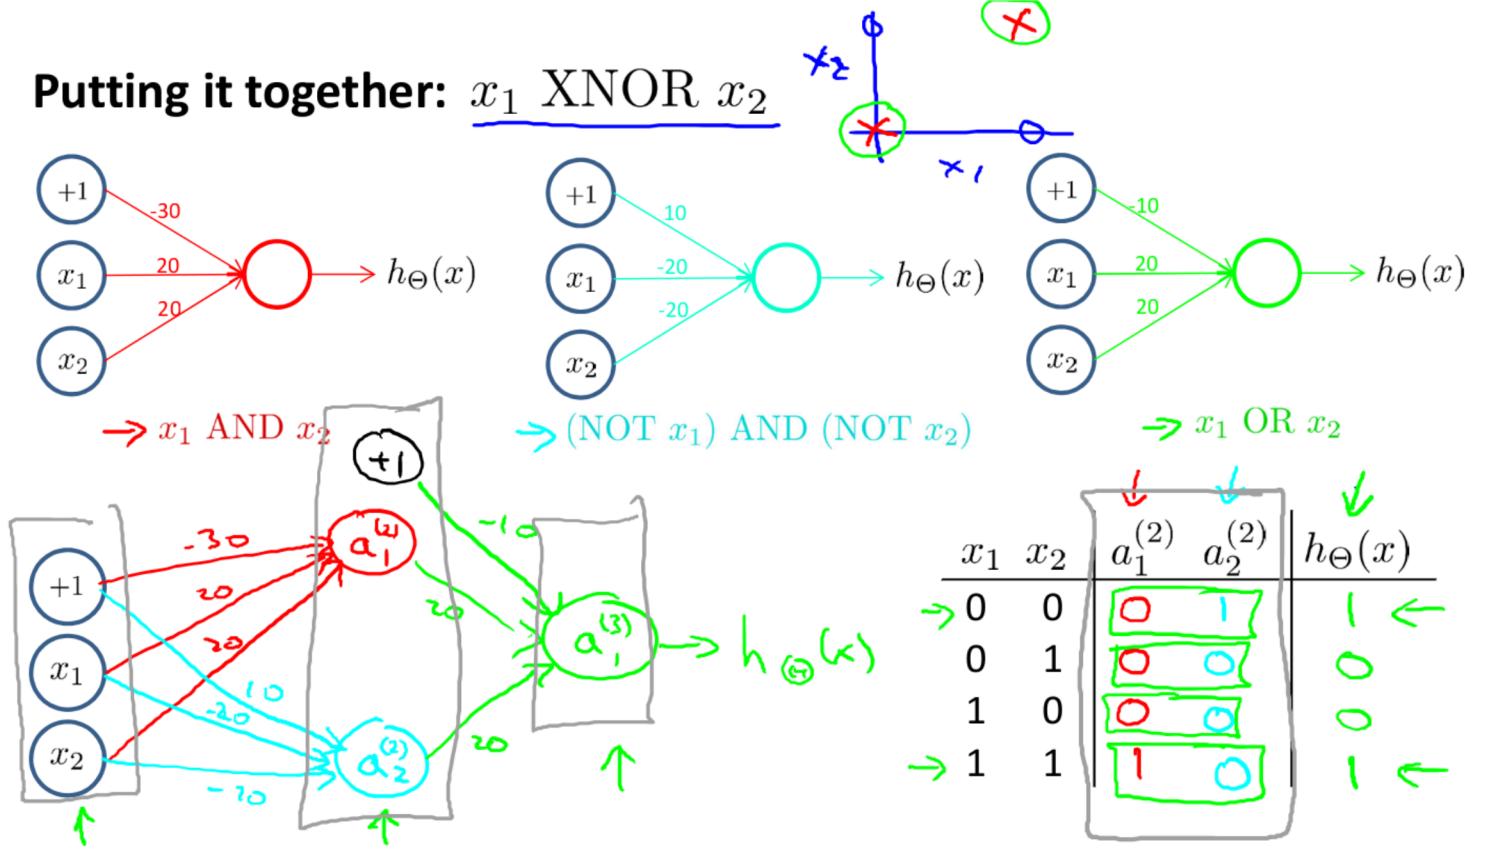
$$h_{\Theta}(x) = g(10 - 20x_1)$$

$\rightarrow (\text{NOT } x_1) \text{ AND } (\text{NOT } x_2)$

$=1$ if and only if
 $\rightarrow x_1 = x_2 = 0$

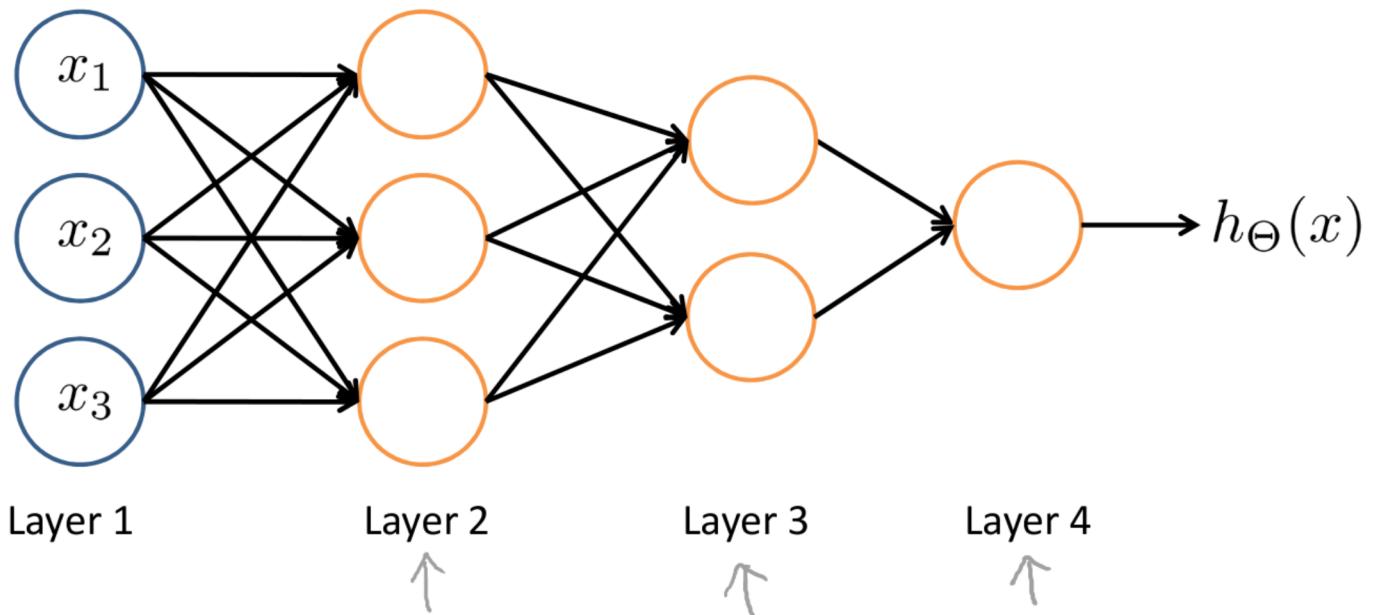
Andrew Ng

Putting it together: x_1 XNOR x_2



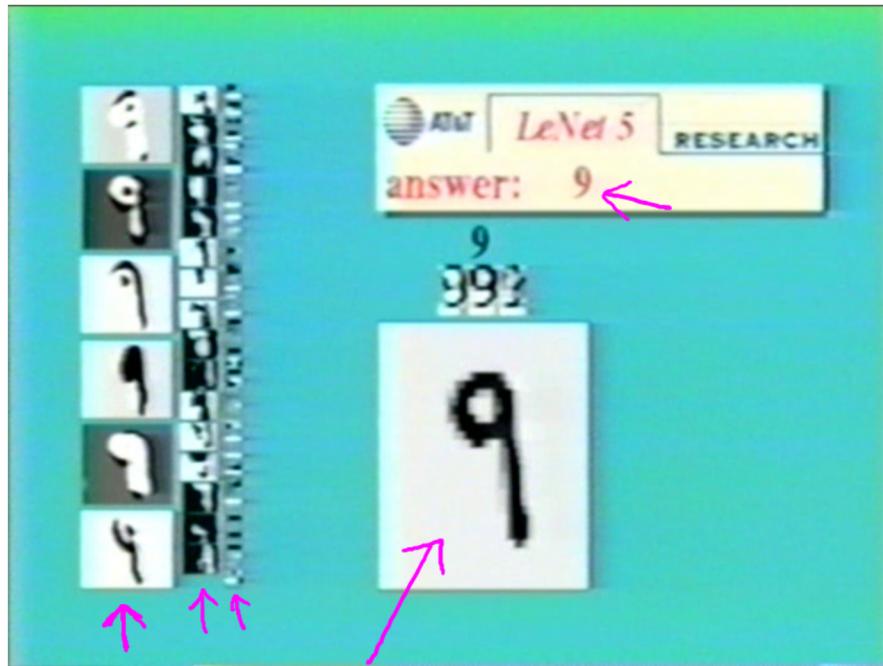
Andrew Ng

Neural Network intuition



Andrew Ng

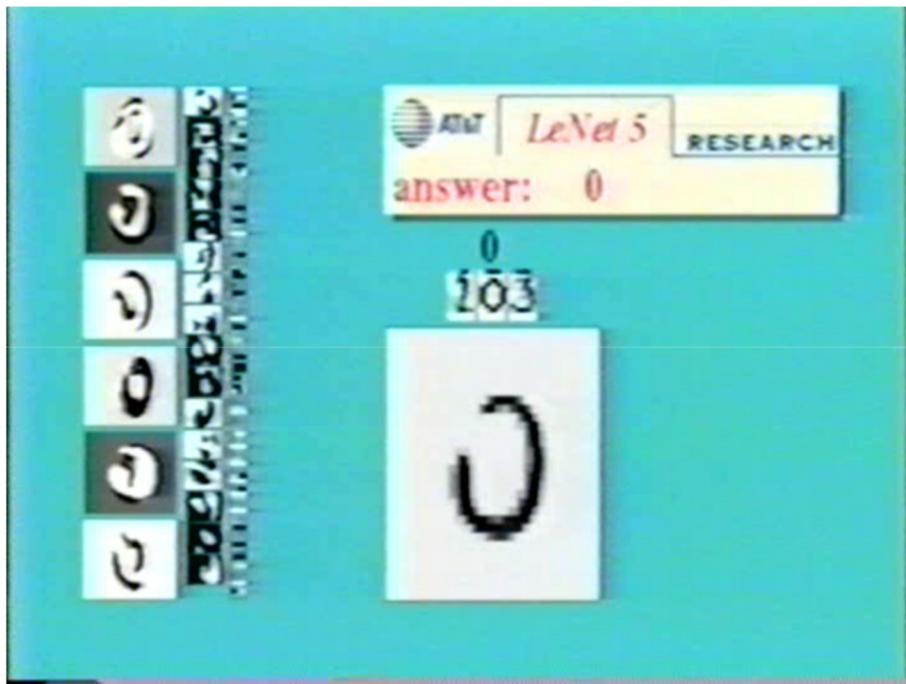
Handwritten digit classification



[Courtesy of Yann LeCun] ↵

Andrew Ng

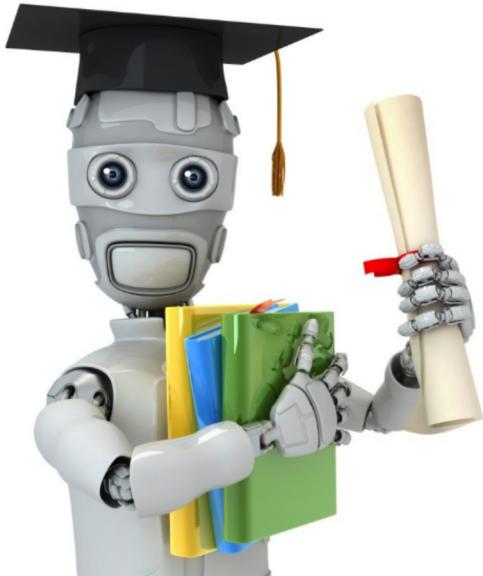
Handwritten digit classification



[Courtesy of Yann LeCun]

Andrew Ng

Andrew Ng



Neural Networks: Representation

Multi-class classification

Machine Learning

Andrew Ng

Multiple output units: One-vs-all.



Pedestrian



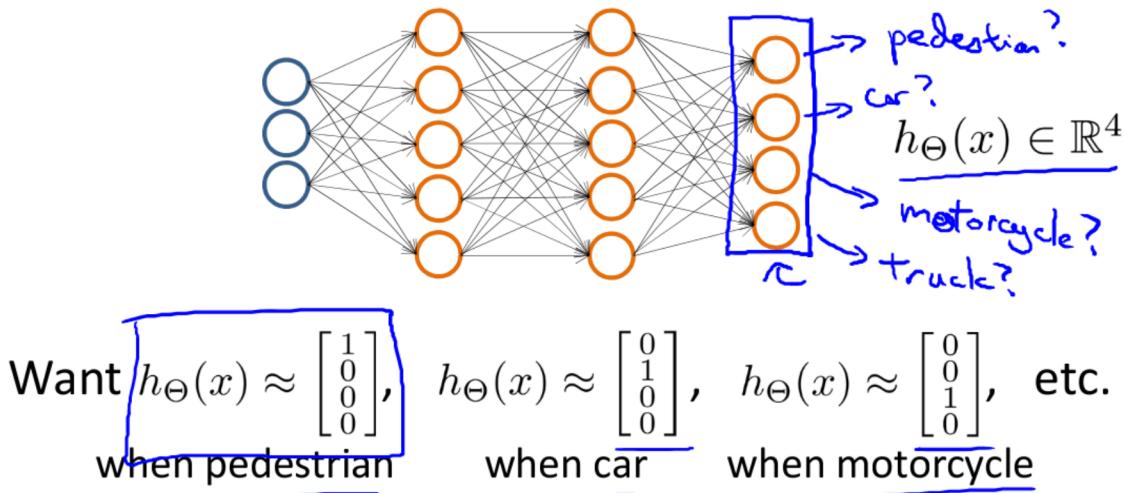
Car



Motorcycle

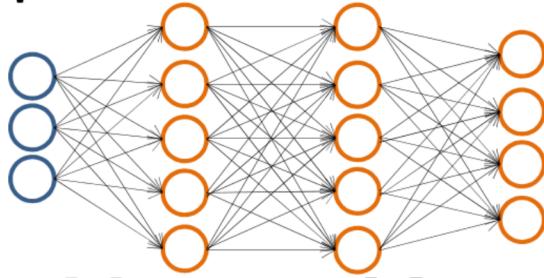


Truck



Andrew Ng

Multiple output units: One-vs-all.



$$h_{\Theta}(x) \in \mathbb{R}^4$$

Want $h_{\Theta}(x) \approx \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}$, $h_{\Theta}(x) \approx \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$, $h_{\Theta}(x) \approx \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$, etc.
 when pedestrian when car when motorcycle

Training set: $(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})$

→ $y^{(i)}$ one of $\begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}$, $\begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$, $\begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$, $\begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$
 pedestrian car motorcycle truck

$$(x^{(i)}, y^{(i)})$$

~~Previously~~
 ~~$y \in \{1, 2, 3, 4\}$~~
 $h_{\Theta}(x^{(i)}) \approx y^{(i)}$
 $\in \mathbb{R}^4$

Andrew Ng

Andrew Ng